

BLG413E-System Programming

Project 1

Merve Elif Demirtaş - 150160706

Gülnur Kaya - 150160154

Buse Dilan Uslan - 150150063

30 October 2018

1 Introduction

In this project, firstly, a new variable -flag- added to the task descriptor of a process and then the behaviour of fork and exit system calls were changed based on the flag and nice value.

2 Changing the Flag Value

Each process has a descriptor related with it. This descriptor includes the information for tracking the process in memory. These information are PID, state, parent process, children, siblings, list of open files, etc. For using the flag value in kernel and also in user space, it has to declared in task descriptor.

- The flag variable declared in task_descriptor. In the task descriptor, the variable added at the end of the file.

```
1464 #if defined(CONFIG_BCACHE) || defined(CONFIG_BCACHE_MODULE)
1465     unsigned int  sequential_io;
1466     unsigned int  sequential_io_avg;
1467 #endif
1468
1469     int myFlag;
1470 };
```

Figure 1: Flag variable declaration in task_struct

- The flag variable added to the task descriptor of mother process at the file where it is initialized in `\include\linux \init_task.h` . The reason of adding the variable to the mother process is fork function. When a fork is called, the child copies all the data of parent.

```

216 .thread_group = LIST_HEAD_INIT(tsk.thread_group), \
217 .thread_node = LIST_HEAD_INIT(init_signals.thread_head), \
218 .myFlag = 0, \
219 INIT_IDS

```

Figure 2: Flag variable initialization in init_task

Under linux-source-3.13.0, set_myFlag/set_myFlag.c function created as below.

```

#include <linux/syscalls.h>
#include <linux/kernel.h>
#include <asm/errno.h>

asmlinkage long sys_set_myFlag(pid_t pid, int flag){

    if((current->cred)->uid == 0 && task_nice > 10){
        struct task_struct *process_ptr = find_task_by_vpid(pid); //return
        the pointer of pid process

        if(process_ptr != NULL){// there is process
            if(flag == 0 || flag == 1){ //valid flag values
                process_ptr-> myFlag = flag;
                return 0;
            }else{
                return EINVAL; //invalid value
            }
        }else{
            return ESRCH; //no such process
        }
    }else{
        return EPERM;
    }
}

```

- In the set_myFlag file, also, a Makefile created and written inside "obj-y := set_myFlag.o" .
- Under linux-source-3.13.0, the Makefile modified adding the set_myFlag to core-y as below.

```

535 # Objects we will link into vmlinux / subdirs we need to visit
536 init-y      := init/
537 drivers-y    := drivers/ sound/ firmware/ ubuntu/
538 net-y       := net/
539 libs-y      := lib/
540 core-y      := usr/ set_myFlag/
541 endif # KBUILD_EXTMOD

```

Figure 3: Makefile of linux-source-3.13.0

- System call table and system call header file was also modified as below.

```

363 354    i386    seccomp          sys_seccomp
364 355    i386    set_myFlag      sys_set_myFlag

```

Figure 4: System call table declaration

```

853 asmlinkage long sys_set_myFlag(pid_t pid, int flag);
854
855 #endif

```

Figure 5: System call header file declaration

3 Changing the Behaviour of Fork and Exit System Call

3.1 Fork System Call

In this part, the behaviour of fork call system changed as when a process has flag value as one, this process can not do fork operation and if it has flag value zero, it can do normal fork operation. When fork system is called, the kernel calls the do_fork function under kernel _fork.c . So, in the do_fork function, copy_process was taken into if function based on flag and nice value as below.

```

1669 if(current->myFlag == 1 && task_nice(current) > 10){
1670     return -ECHILD;
1671 }else{
1672     p = copy_process(clone_flags, stack_start, stack_size, child_tidptr, NULL, trace);
1673 }

```

Figure 6: do_fork function

This modification tested as below.

```
#include <stdio.h>
#include <stdlib.h>
#include <asm/errno.h>
#include <unistd.h>
#include <string.h>

#define set_myFlag 355

int main(){

    int flag = 1;

    int check_return;
    printf("getpid(): %d , getppid(): %d \n", getpid(), getppid());
    check_return = syscall(set_myFlag, getpid(), flag); //set flag value
        to mother process

    printf("Return value of set_myFlag: %s\n", strerror(check_return));

    int f;
    f = fork();
    printf("f value: %d\n", f);

    if(f == 0){ //child process
        printf("Child pid: %d child parent pid:%d\n", getpid(), getppid());

        return 0;
    }else if(f < 0){
        printf("Return value: %s\n", strerror(f));
        printf("getpid(): %d , getppid(): %d \n", getpid(), getppid());
    }else{
        printf("Parent pid: %d\n", getpid());
    }

    return 0;
}
```

3.2 Exit System Call

In this part, the behaviour of exit system call was changed as when a child process calls the exit function, it will exit with killing all the siblings. The Linux kernel uses a circular doubly-linked list of struct task _structs to store process descriptors.




Figure 7: Circular doubly-linked list of struct task _structs

So, the sibling list of child iterated over list and killed with sys_kill function as below.

```
704 void do_exit(long code)
705 {
706     struct task_struct *tsk = current;
707     int group_dead;
708
709     struct list_head *head_of_list;
710     struct task_struct *task;
711
712     if(tsk->myFlag == 1 && task_nice(tsk) > 10){
713         list_for_each(head_of_list, &tsk->sibling) { //iterate over sibling list
714             task = list_entry(head_of_list, struct task_struct, sibling);
715             sys_kill(task->pid, SIGKILL);
716         }
717     }
718 }
```

Figure 8: do_exit function under kernel_exit.c

This modification tested as below.

```
#include <stdio.h>
#include <stdlib.h>
#include <asm/errno.h>
#include <unistd.h>
#include <string.h>

#define set_myFlag 355

int main(){

    int f, i, check_return, flag, tmp;
```

```

printf("FIRST -> getpid(): %d, getppid(): %d \n", getpid(),
      getppid());

for(i = 0; i < 5; i++){
    f = fork();

    if(i == 4){
        if(f == 0){//child process
            flag = 1;

            int check_return;
            printf("PROCESS CALLS THE EXIT -> getpid(): %d ,
                  getppid(): %d \n", getpid(), getppid());
            check_return = syscall(set_myFlag, getpid(), flag);//try
                           to set flag to child process

            printf("PROCESS CALLS THE EXIT -> Return value of
                  set_myFlag: %s\n", strerror(check_return));

            sleep(5);
            printf("The child that calls the exit call was died!\n");
            exit(0);
        }
    }

    if(f == 0){
        printf("Waiting process child: getpid(): %d , getppid(): %d
              \n", getpid(), getppid());
        while(1);
    }
}

sleep(10);
return 0;
}

```
